

Database Management Systems

INFO 210

Entity Relationships

Lecture 3

Franz Wotawa

TU Graz, Institut for Software Technologie

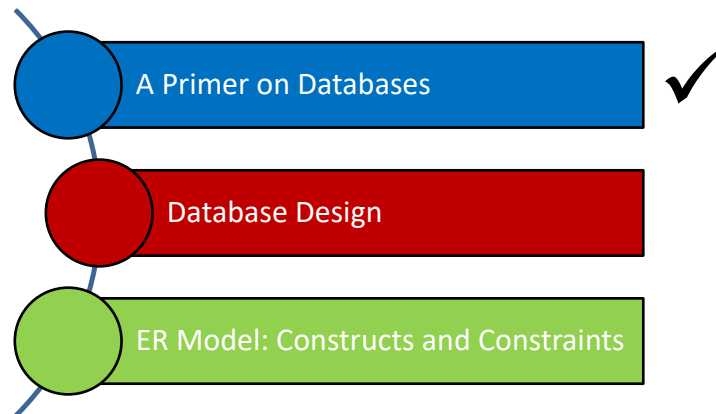
Inffeldgasse 16b/2

wotawa@ist.tugraz.at

Today...

- **Last Session:**
 - Intro do databases
 - Conceptual Modeling
- **Today's Session:**
 - Introduction to databases and database systems (*recap*)
 - Main steps involved in designing databases
 - Constructs of the entity relationship (ER) model
 - Integrity constrains that can be expressed in the ER model

Outline



Some Definitions

- A **database** is a collection of data which describes one or many real-world enterprises
 - E.g., a university database might contain information about **entities** like students and courses, and **relationships** like a student enrollment in a course
- A **DBMS** is a software package designed to store and manage databases
 - E.g., DB2, Oracle, MS SQL Server, MySQL and **Postgres**
- A **database system** = (Big) Data + DBMS + Application Programs

Data Models

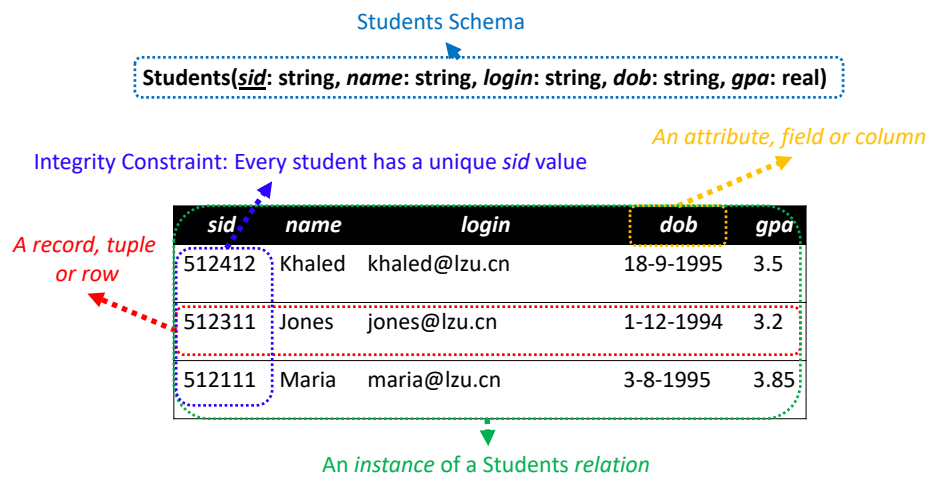
- The user of a DBMS is ultimately concerned with some real-world enterprises (e.g., a University)
- The data to be stored and managed by a DBMS *describes* various aspects of the enterprises
 - E.g., The data in a university database describes students, faculty and courses entities and the relationships among them
- A **data model** is a collection of high-level data description constructs that hide many low-level storage details
- A widely used data model called the **entity-relationship (ER) model** allows users to pictorially denote entities and the relationships among them

The Relational Model

- The **relational model** of data is one of the most widely used models today
- The central data description construct in the relational model is the **relation**
- A relation is basically a **table** (or a **set**) with **rows** (or **records** or **tuples**) and **columns** (or **fields** or **attributes**)
- Every relation has a **schema**, which describes the columns of a relation
- Conditions that records in a relation must satisfy can be specified
 - These are referred to as **integrity constraints**

The Relational Model: An Example

- Let us consider the student entity in a university database




Queries in a DBMS

- The “ease” (plus efficiency & effectiveness) with which information can be queried from a database determines its value to users
- A DBMS provides a specialized language, called the **query language**, in which queries can be posed
- The relational model supports powerful query languages
 - Relational calculus**: a formal language based on mathematical logic
 - Relational algebra**: a formal language based on a collection of operators (e.g., selection and projection) for manipulating relations
 - Structured Query Language (SQL)**:
 - Builds upon relational calculus and algebra
 - Allows creating, manipulating and querying relational databases
 - Can be embedded within a host language (e.g., Java)

Concurrent Execution and Transactions

- An important task of a DBMS is to *schedule* concurrent accesses to data so as to improve performance

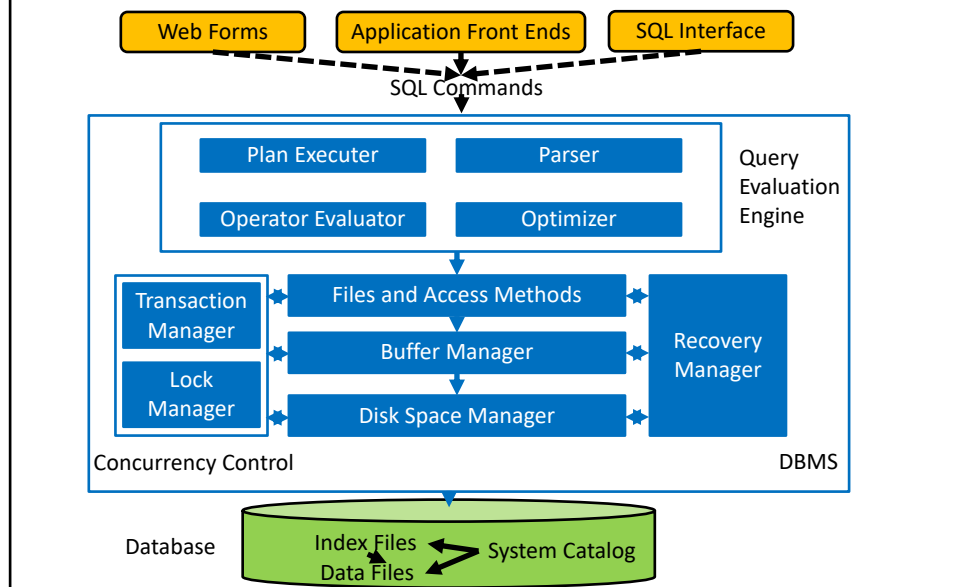


- When several users access a database *concurrently*, the DBMS must order their requests carefully to avoid conflicts
 - E.g., A check might be cleared while account balance is being computed!
- DBMS ensures that conflicts do not arise via using a **locking protocol**
 - Shared vs. Exclusive locks

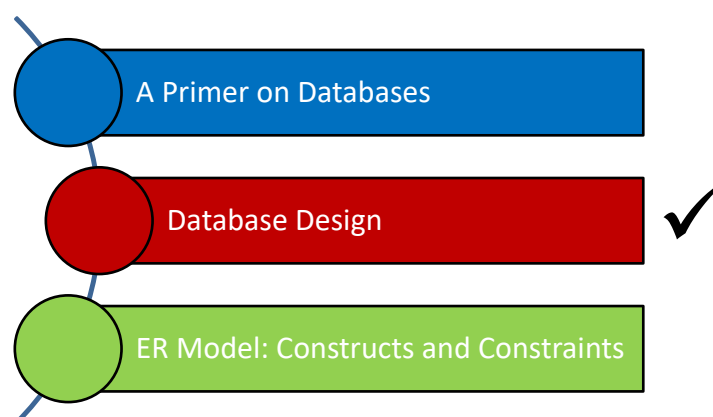
Ensuring Atomicity

- Transactions can be interrupted before running to completion for a variety of reasons (e.g., due to a system crash)
- DBMS ensures atomicity (**all-or-nothing property**) even if a crash occurs in the middle of a transaction
- This is achieved via maintaining a **log** (i.e., history) of all writes to the database
 - *Before* a change is made to the database, the corresponding log entry is forced to a safe location (this protocol is called **Write-Ahead Log** or **WAL**)
 - After a crash, the effects of partially executed transactions are *undone* using the log

The Architecture of a Relational DBMS



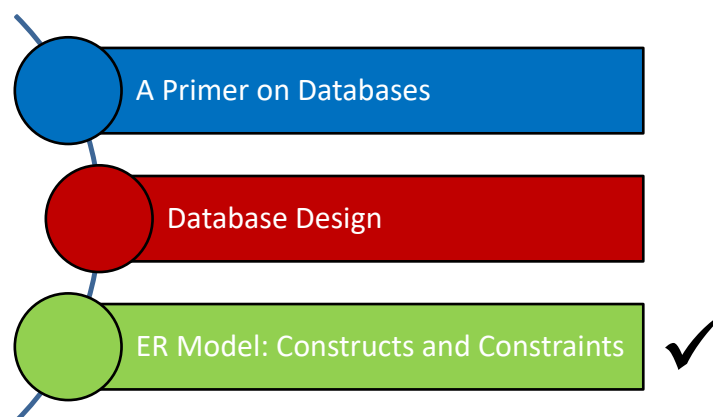
Outline



Database Design

- **Requirements Analysis**
 - Users needs
- **Conceptual Design**
 - A high-level description of the data (e.g., using the ER model)
- **Logical Design**
 - The conversion of an ER design into a relational database schema
- **Schema Refinement**
 - Normalization (i.e., restructuring tables to ensure some desirable properties)
- **Physical Design**
 - Building indexes and clustering some tables
- **Security Design**
 - Access controls

Outline



Entities and Entity Sets


Entity:

- A real-world object distinguishable from other objects in an enterprise (e.g., University, Students and Faculty)
- Described using a set of *attributes*

Entity set:

- A collection of similar entities (e.g., *all* employees)
- All entities in an entity set have the same set of attributes (until we consider *ISA* hierarchies, anyway!)
- Each entity set has a *key*
- Each attribute has a *domain*

Tools and An ER Diagram




Entities ('Entity Sets')



Attributes

"ssn" is the
primary key,
hence, underlined



Relationship and Relationship Sets


Relationship:

- Association among two or more entities (e.g., Rui is teaching INFO210)
- Described using a set of attributes


Relationship set:

- Collection of similar relationships
- Same entity set could participate in different relationship sets, or in different “roles” in the same set


More Tools and ER Diagrams



Relationships ('rel. sets')
and mapping constraints




A Binary Relationship



A Self-Relationship

Ternary Relationships


- Suppose that departments have offices at different locations and we want to record the locations at which each employee works
- Consequently, we must record an association between an employee, a department and a location



This is referred to as a "Ternary Relationship" (vs. Self & Binary Relationships)

Key Constraints


- Consider the "Employees" and "Departments" entity sets with a "Manages" relationship set
 - An employee can work in *many* departments
 - A department can have *many* employees
 - Each** This restriction is an example of a "key constraint"



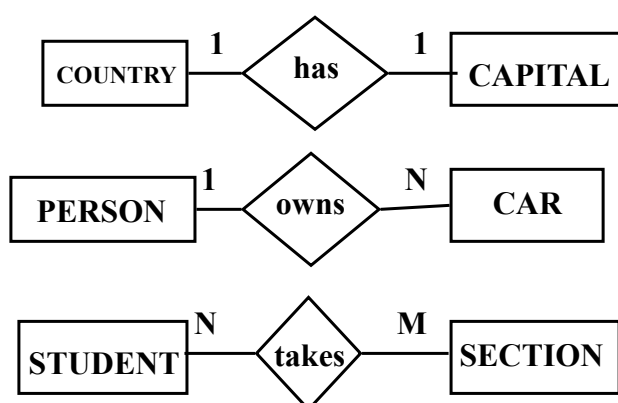
Key constraints are denoted by *thin arrows*

Cardinalities

- Entities can be related to one another as “one-to-one”, “one-to-many” and “many-to-many”
 - This is said to be the **cardinality** of a given entity in relation to another

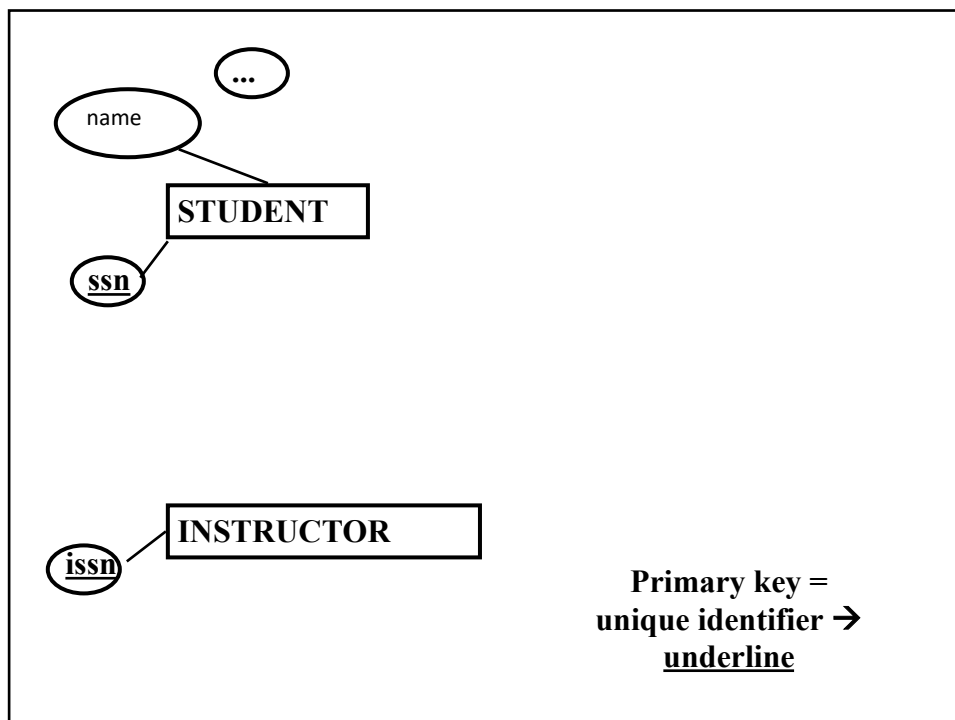


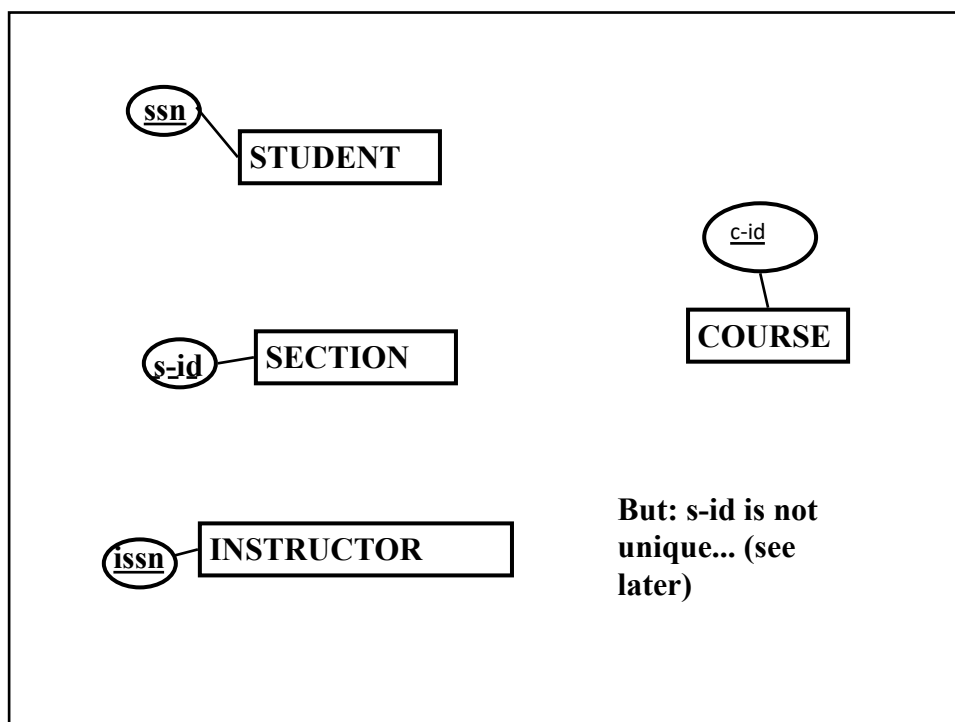
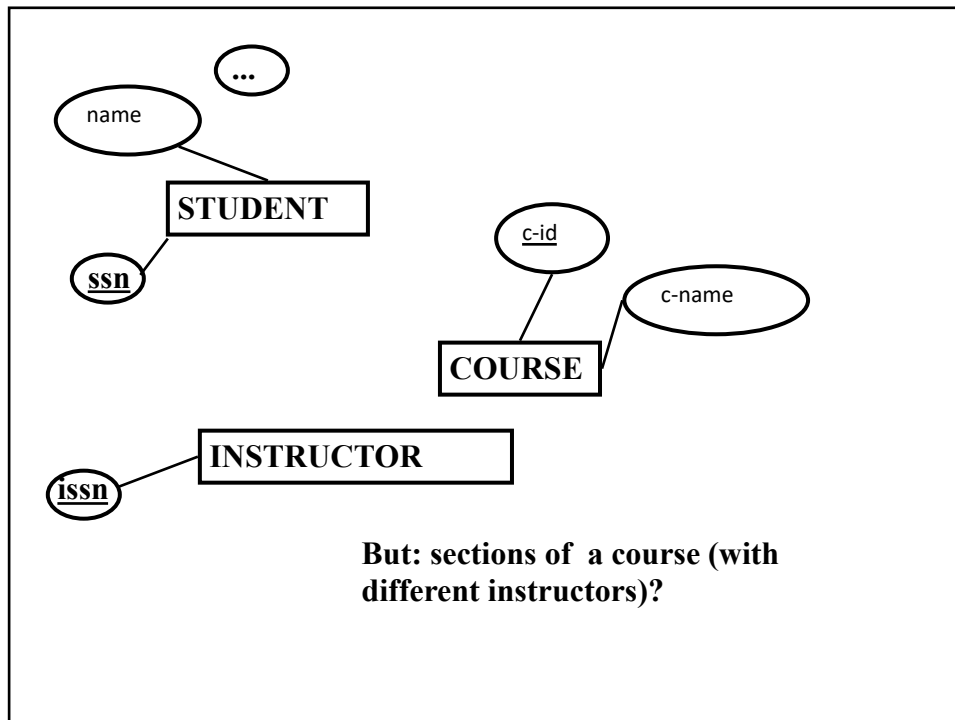
Cardinalities: Examples

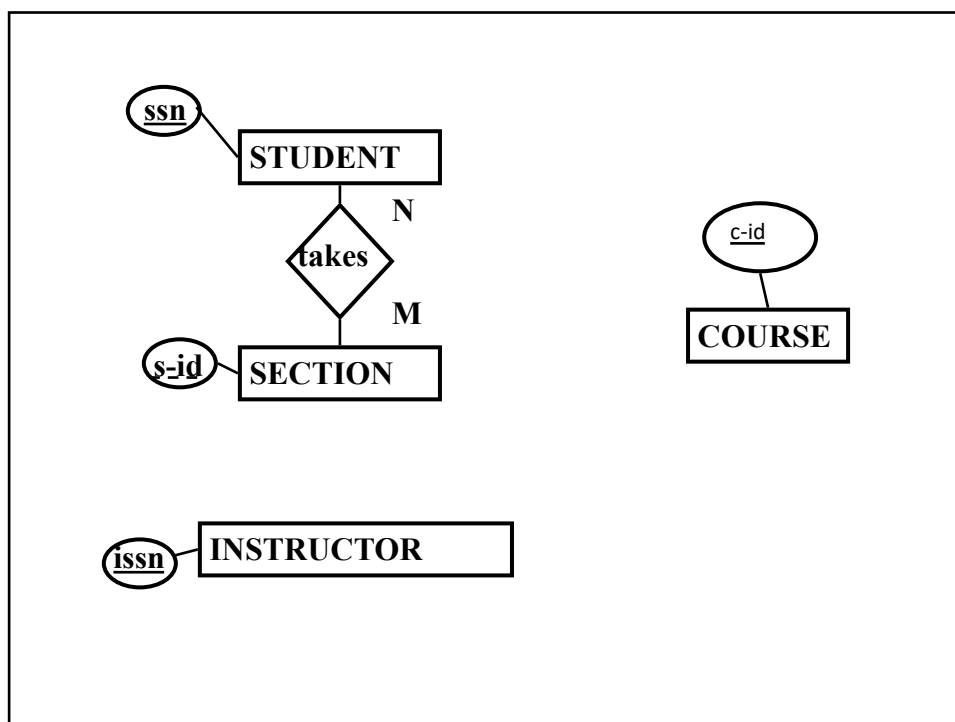
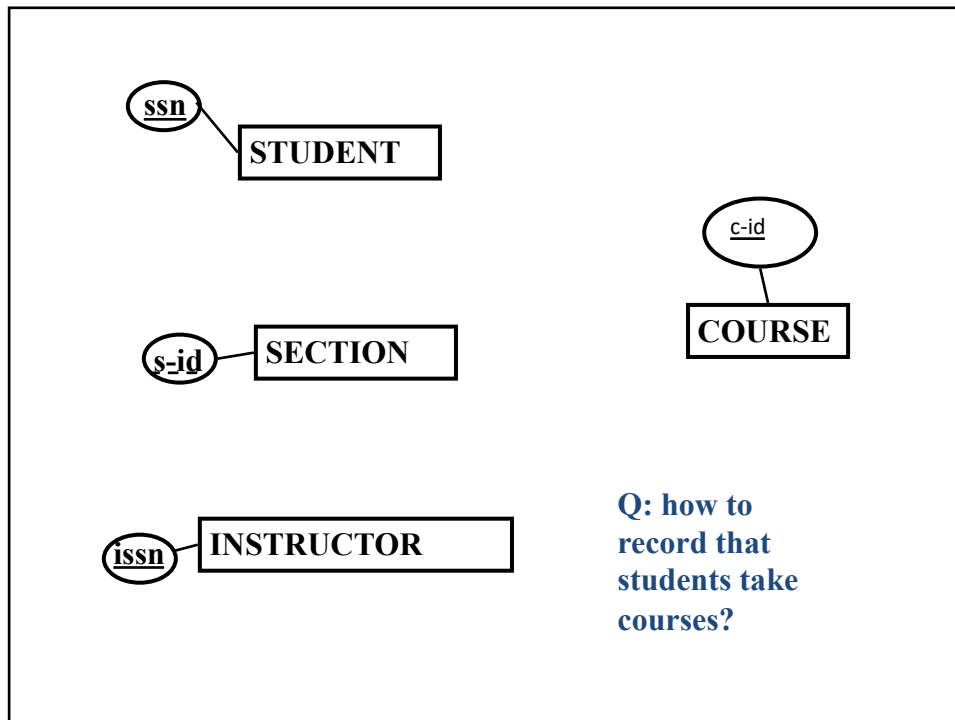


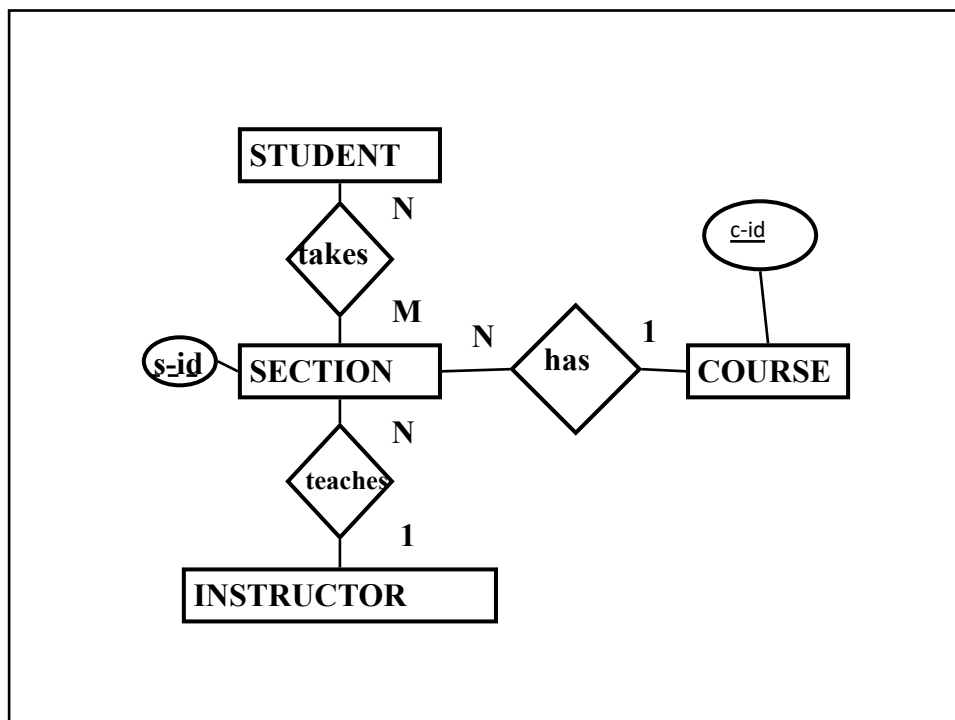
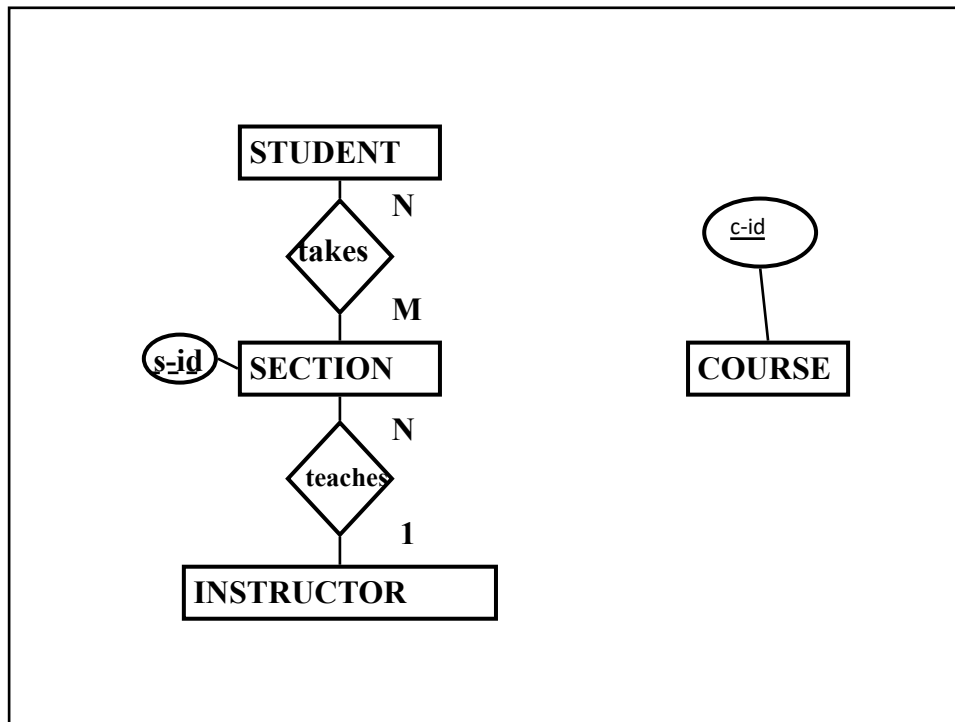
A Working Example

- **Requirements:** Students take courses offered by instructors; a course may have multiple sections; one instructor per section
- **How to start?**
 - Nouns -> entity sets
 - Verbs -> relationship sets











Participation Constraints

- Consider again the “Employees” and “Departments” entity sets as well as the “Manages” relationship set
 - Should every department have a manager?
 - If so, this is a **participation constraint**
 - Such a constraint entails that every Departments entity must appear in an instance of the Manages relationship
 - The participation of Departments in Manages is said to be **total** (vs. **partial**)



Total vs. Partial Participations

Total, Total



??



??




Total vs. Partial Participations

Total, Total



Partial, Total




??




Total vs. Partial Participations

Total, Total



Partial, Total



Partial, Total




Weak Entities

- A **weak entity** can be identified uniquely only by considering the primary key of another (*owner*) entity
 - Owner entity set and weak entity set must participate in a one-to-many relationship set (one owner, many weak entities)
 - Weak entity set must have total participation in this **identifying relationship set**
- The set of attributes of a weak entity set that uniquely identifies a weak entity for a given owner entity is called **partial key**

Weak Entities: An Example

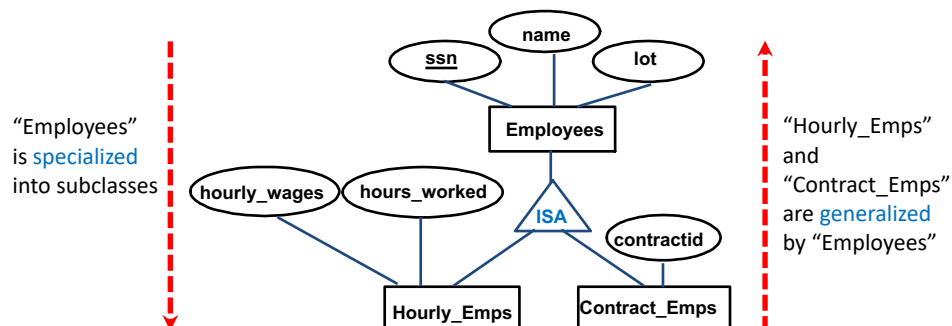
- “Dependents” has no unique key of its own
 - “Dependents” is a weak entity with partial key “pname”
 - “Policy” is an identifying relationship set
 - “pname” + “ssn” are the primary key of “Dependents”



Weak entities and identifying relationships are drawn using *thick lines*


ISA ('is a') Hierarchies

- Entities in an entity set can sometimes be classified into subclasses (this is “kind of similar” to OOP languages)
- If we declare B **ISA** A, every B entity is also considered to be an A entity



Overlap and Covering Constraints


- Overlap constraints**
 - Can an entity belong to both 'B' and 'C'?
- Covering constraints**
 - Can an 'A' entity belong to neither 'B' nor 'C'?



Overlap Constraints: Examples

Overlap constraints


- Can John be in Hourly_Emps and Contract_Emps? Intuitively, *no*
- Can John be in Contract_Emps and in Senior_Emps? Intuitively, *yes* →
"Contract_Emps OVERLAPS Senior_Emps"



Covering Constraints: Examples

Covering constraints

- Does every one in Employees belong to a one of its subclasses? Intuitively, *no*
- Does every Motor_Vehicles entity have to be either a Motorboats entity or a Cars entity? Intuitively, *yes* →
"Motorboats AND Cars COVER Motor_Vehicles"

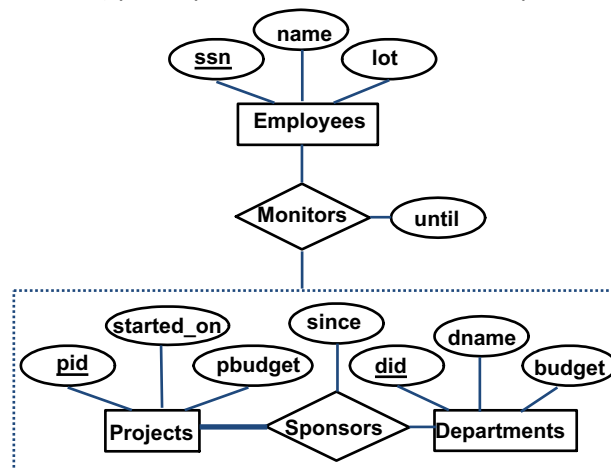


More Details on ISA Hierarchies

- Attributes are *inherited* (i.e., if B **ISA** A, the attributes defined for a B entity are the attributes for A *plus* B)
- We can have **many** levels of an ISA hierarchy
- Reasons for using ISA:
 - To add descriptive attributes specific to a subclass
 - To identify entities that participate in a relationship

Aggregation

- Aggregation allows indicating that a relationship set (identified through a *dashed box*) participates in another relationship set



Next Class

Continue the ER Model and
Start with the relational Model