

# **Database Management Systems INFO 210**

**NoSQL – Recent developments in database  
research and application  
Lecture 14**

**Franz Wotawa**  
TU Graz, Institut for Software Technologie  
Inffeldgasse 16b/2  
wotawa@ist.tugraz.at

## **Notes**

This slides are provided from Prof. Dr. Reinhard Pichler from the Institute of Logic and Computation, Databases & Artificial Intelligence Group, TU Vienna.

The content is not completely representing all aspects of today's database research and applications but at least the main streams

## References

### ❖ Textbooks:

- Pramod J. Sadalage and Martin Fowler: NoSQL Distilled. Addison-Wesley 2013
- Lena Wiese: Advanced Data Management for SQL, NoSQL, Cloud and Distributed Databases. de Gruyter, 2015

### ❖ Scientific Articles:

- Felix Gessert, Wolfram Wingerath, Steffen Friedrich, Norbert Ritter: NoSQL database systems: a survey and decision guidance. Computer Science - R&D 32(3-4): 353-365 (2017).

## References (cont.)

### ❖ Further Scientific Articles:

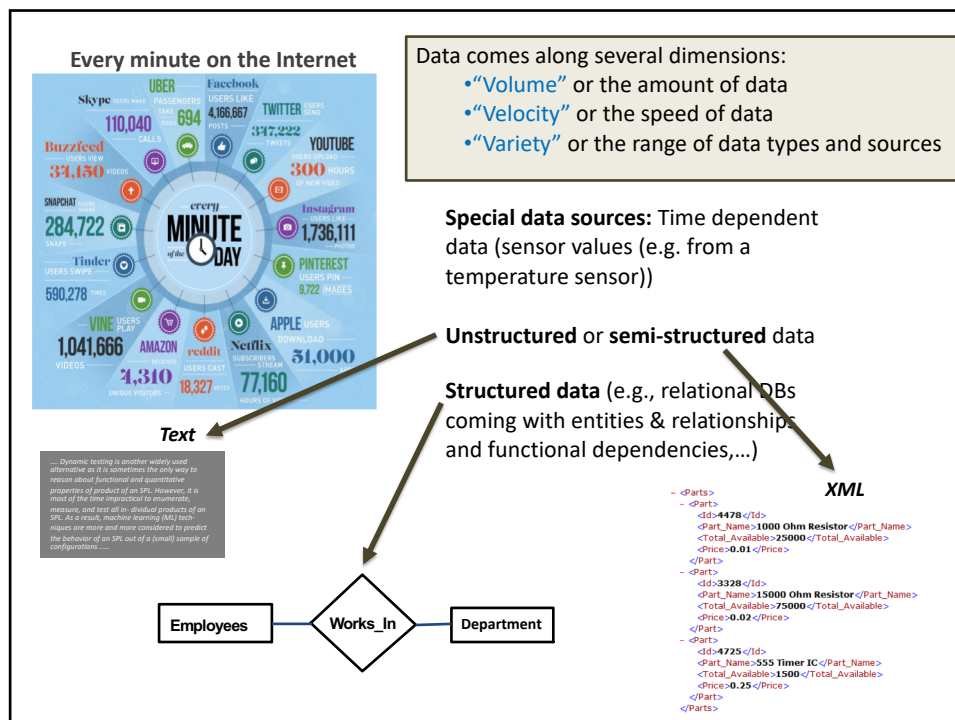
- Eric A. Brewer: CAP twelve years later: How the "rules" have changed. IEEE Computer 45(2): 23-29 (2012)
- Daniel Abadi: Consistency Tradeoffs in Modern Distributed Database System Design: CAP is Only Part of the Story. IEEE Computer 45(2): 37-42 (2012)
- Dan Pritchett: BASE: An ACID Alternative. ACM Queue 6(3): 48-55 (2008)

### ❖ Copyright of all figures is with one of these resources.

## Contents

- ❖ **Motivation**
- ❖ Four Categories of NoSQL Systems
- ❖ Further Aspects of Data Modelling
- ❖ Document Stores
- ❖ MongoDB Primer

5



## Relational DBMSs

### ❖ Strengths:

- Strong formal foundation (logic, relational algebra)
- Standard data model and query language (SQL)
- Concurrency: support multi-user access; transactions
- Data integration: store data of many applications in one database

### ❖ Previous alternative proposals:

- XML Stores, Object-Oriented DBMSs
- Have been integrated into RDBMSs and complement/extend now their functionality

7

## ACID property of Relational DBMS

- Transactions have the following four properties:
  - **Atomicity** – Either a transaction completes or nothing happens at all.
  - **Consistency** – A transaction must start in a consistent state and leave the system in a consistent state.
  - **Isolation** – Intermediate results of a transaction are not visible outside the current transaction.
  - **Durability** – Once a transaction was committed, the effects are persistent, even after a system failure.

## Limitations of Relational DBMSs

- ❖ **Schema**: must be known before any data can be added
- ❖ **Data structure**:
  - RDBMS are designed for structured data; limited support for unstructured data (text, pictures/videos)
  - Requires horizontal and vertical homogeneity of data (same structure per row, same type per column).
  - Data belonging to an entity is split into several tables
  - **Impedance Mismatch**: relational model vs. in-memory data structures (e.g.: OO, nested structures)
- ❖ **Cost**:
  - major vendors: closed source with strict licensing fees
  - usual licensing model: pay per user / computing node

9

## Time Consumption in RDBMSs

- ❖ M. Stonebraker et al. have found that in a typical RDBMS, only 6.8% of the execution time is spent on "useful work".
- ❖ **The rest** is spent on:
  - buffer management (34.6%): disk I/O
  - locking (16.3%): to control concurrency
  - latching (14.2%): to protect shared data structures from race conditions caused by multi-threading
  - logging (11.9%): to guarantee durability
  - other code parts (16.2%)
- ❖ **Conclusion**: the end of the "one-size-fits-all" era

Reference: Michael Stonebraker, et al.: The End of an Architectural Era (It's Time for a Complete Rewrite). VLDB 2007: 1150-1160

10

## New Data Management Challenges

- ❖ **Volume:** processing "big data"
- ❖ **Complex data structures:**
  - Nested structures
  - Complex relationships (graph structure)
- ❖ **Schema Independence / Schema Evolution:**
  - to facilitate integration of data from various sources
  - coping with changing structure of the data
  - allow self-descriptiveness
- ❖ **Sparseness:**
  - Many data items are unknown or non-existent
  - Would result in many NULLs in relational tables

11

## Reasons for Emergence of New DB Systems

- ❖ **Application vs. integration database:**
  - Shift from idea of integration DB to application DB, i.e.: each application maintains its own database
  - Data exchange via web services
  - Use rich structure for communication (to reduce number of round trips in the interaction)
  - advantages of application DB: more freedom in data modelling, shift features to application (e.g. security)
- ❖ **Use of Clusters**
  - Scale out to cope with huge amounts of data;
  - RDBMSs are not well-suited for running on clusters.

12

## The Term "NoSQL"

### ❖ History:

- June 2009: meetup in San Francisco for "**open-source, distributed, non-relational databases**".
- organized at the time of the Hadoop Summit 2009;
- search for a Twitter hashtag without many Google hits;
- **NoSQL** was chosen

### ❖ Some of the **Systems presented at this 2009 meetup**:

<https://blog.oskarsson.nu/post/22996140866/nosql-debrief>

- Voldemort, Cassandra, Dynomite, HBase, Hypertable, CouchDB, **MongoDB**

13

## The Term "NoSQL" (cont.)

### ❖ **NoSQL = "Not only SQL"**

### ❖ NoSQL as an umbrella term for all those DB systems that share **most of the following properties**:

- open-source, distributed, non-relational
- schemaless
- restricted querying capabilities (e.g., no joins)
- horizontally scalable; running well on clusters
- no (or weaker form of) transactions, not guaranteeing ACID properties, use BASE instead

**Basically Available, Soft state, Eventual consistency**

14

## Pioneering Contributions by Companies

### ❖ Early systems:

- Google: Bigtable (column family store)
- Amazon: DynamoDB (key-value store)
- Facebook: Cassandra (column family store), later used by Twitter
- LinkedIn: Project Voldemort (open-source implementation of Amazon's DynamoDB)

### ❖ Characteristics of early systems:

- Huge number of users
- Huge number of reads/writes
- Non-critical transactional systems (ACID not needed)

15

## Benefits of NoSQL Database Systems

### ❖ High performance:

- Handling huge amounts of data
- Specific optimizations (e.g.: compression in column stores; path queries in graph databases)

### ❖ Horizontal scalability:

- Add nodes to the cluster as load or DB grows

### ❖ Schemaless:

- Simpler storage mechanisms: performance
- No schema needed upfront, flexible when data changes

### ❖ No (or weaker forms of) transactions:

- Simpler transaction handling; not blocking
- Reduced logging requirements

16



## Disadvantages of NoSQL Systems

- ❖ **No standardized Query Language:**
  - Users need to learn a separate query language for each system (even within the same category)
- ❖ **Weaker form (or no support of) transactions**
  - no ACID support (in particular, consistency)
- ❖ **Limited support of ensuring integrity:**
  - Responsibility for uniqueness constraints, referential integrity, etc. is shifted to the application.

17

## Contents

- ❖ Motivation
- ❖ **Four Categories of NoSQL Systems**
- ❖ Further Aspects of Data Modelling
- ❖ Document Stores
- ❖ MongoDB Primer

18

## Four Categories

- ❖ **Key-value stores:**
  - Stores values indexed by keys
- ❖ **Document stores:**
  - Similar to key-value stores but DBMS can understand the format of the value; typically JSON, XML
- ❖ **Columnar databases, column family stores:**
  - Stores data organized in columns
  - Can be thought of as two-dimensional indexing of values: first by id (like key) then by column name
- ❖ **Graph Databases:**
  - Uses graphs to represent data; vertices represent entities; edges represent (directed) relationships

19

## Examples of Systems in the Four Categories

See <http://nosql-database.org/> (over 200 systems);  
(see also <https://db-engines.com/en/ranking>)

- ❖ **Key-value stores:**
  - Riak, Redis, BerkeleyDB, DynamoDB (Amazon); open source versions: Dynamite, Project Voldemort
- ❖ **Document stores:**
  - **MongoDB**, CouchDB, Terrastore
- ❖ **Columnar databases, column family stores:**
  - MonetDB, HBase, Cassandra, Hypertable, Accumulo
- ❖ **Graph Databases:**
  - Neo4J, HyperGraphDB, Infinite Graph

20

## Complex Structures

### ❖ "Aggregate Orientation":

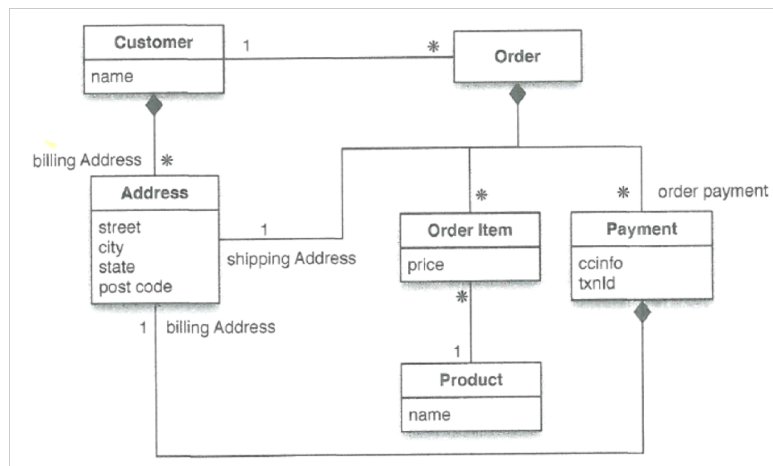
- Data organization in RDBMS: by tuples
- many applications: access to more complex structures
- **aggregate**: = collection of objects that we want to treat as a unit, e.g. unit of access to storage, unit of updates, unit of communication
- in data modelling: **denormalization**
- main goal of NoSQL data modelling: read optimization

### ❖ Consequences of aggregate orientation:

- Well-suited for running on a cluster
- Consequence for transactions: easy to support atomic manipulation of single aggregate; difficult to support atomic manipulation of several aggregates

21

## Complex Structures: Example



22

## Complex Structures: Example (cont.)

```
// in customers
{
  "id":1,
  "name":"Martin",
  "billingAddress":[{"city":"Chicago"}]
}

// in orders
{
  "id":99,
  "customerId":1,
  "orderItems":[
    {
      "productId":27,
      "price": 32.45,
      "productName": "NoSQL Distilled"
    }
  ],
  "shippingAddress":[{"city":"Chicago"}]
```

```
// in customers
{
  "customer": {
    "id": 1,
    "name": "Martin",
    "billingAddress": [{"city": "Chicago"}],
    "orders": [
      {
        "id":99,
        "customerId":1,
        "orderItems":[
          {
            "productId":27,
            "price": 32.45,
            "productName": "NoSQL Distilled"
          }
        ],
        "shippingAddress":[{"city":"Chicago"}]
```

23

## Aggregates in NoSQL Systems

- ❖ **Key Value Stores:**
  - arbitrary structure of values: opaque to DB system; only known to the application.
- ❖ **Document Stores:**
  - nested structure (XML, JSON) visible to the DB system
- ❖ **Column-Oriented DBs:**
  - Organize columns that are usually accessed together in column families (pioneer: Bigtable)
- ❖ **Graph Databases:**
  - aggregate-ignorant (like RDBMSs)
  - little information stored per entity; most important information lies in the relationships.

24

## Contents

- ❖ Motivation
- ❖ Four Categories of NoSQL Systems
- ❖ **Further Aspects of Data Modelling**
- ❖ Document Stores
- ❖ MongoDB Primer

25

## Relationships

- ❖ **Data accesses:**
  - aggregates help to put together data that is commonly accessed together.
  - Relationships between aggregates: usually, not all related data is combined in one aggregate; in particular, applications may have different aggregate boundaries.
- ❖ **Example:**
  - Customers vs. orders: one application may access customer data together with order information; other application may access only the orders.
  - Shall we store customers and orders as one aggregate or as separate aggregates?

26

## Relationships (cont.)

- ❖ **Modelling inter-aggregate relationships:**
  - IDs in one aggregate to reference other aggregates
  - In general, this reference-semantics is not visible to the DB system; however, e.g. Riak allows link information in metadata and supports link-walking
  - Graph DBs are optimized for traversing relationships.
- ❖ **Challenge with inter-aggregate relationships:**
  - how to ensure atomicity of changes to related aggregates (ACID property of transactions in RDBMS)?

27

## Materialized Views

- ❖ **Materialized Views in RDBMSs:**
  - Lack of aggregate structure in RDBMS: easier to support different ways of accessing the data
  - Views: virtual tables; convenient for encapsulation
  - Materialized views: views cached on disk
- ❖ **NoSQL Systems:**
  - Typically don't have views but some allow precomputed and cached queries (referred to as materialized views)
- ❖ **Update strategies:**
  - Eagerly: whenever base data is updated
  - Lazily: update materialized views at regular intervals

28

## Schemaless Databases

### ❖ Freedom and flexibility in NoSQL Systems:

- key-value store: arbitrary contents in value
- document store: arbitrary structure of documents
- Column-oriented: arbitrary combination of columns
- flexibility mainly applies to data within an aggregate, not when aggregate boundaries have to be modified.

### ❖ "Schemaless":

- schema on write (RDBMS) vs. schema on read
- implicit schema: DB system is ignorant of schema but application programs assume a schema
- may be difficult to maintain: requires understanding of the schema in the application code.

29

## Contents

- ❖ Motivation
- ❖ Four Categories of NoSQL Systems
- ❖ Further Aspects of Data Modelling
- ❖ **Document Stores**
- ❖ MongoDB Primer

30

## Document Stores

### ❖ Document stores:

- Similar to key-value stores but DBMS can understand the format of the value
- Typical data formats: XML, JSON (JavaScript Object Notation), BSON (binary JSON)
- Allows advanced querying (e.g.: filters on field values) and relationships (references) between data objects
- Joins usually not supported for performance reasons; consequence: denormalizing data when modelling or combining documents on application level.
- Atomicity usually (only) of documents guaranteed
- **In this lecture:** look at MongoDB

31

## Introduction to MongoDB

### ❖ Open Source

### ❖ Data format: stores data in BSON

### ❖ Name: from "humongous ": = enormous, gigantic.

### ❖ Brief History:

2007: started as part of developing cloud computing stack

2009: version 1 of MongoDB released

Constantly extended feature set, e.g.:

- Sharding (v1.6, 2010), Journaling (v1.8, 2011)
- Aggregation framework (v2.2, 2012)
- Document validation (v3.2, 2015)
- Multi-document ACID transactions (v.4.0, 2018)

32



## Typical Use Cases

- ❖ **Log data:** rather than storing pure text, extract the various fields from a log file (e.g., host, user, time, path, etc.)
- ❖ **e-commerce, e.g.:**
  - product catalog
  - shopping cart
- ❖ **Inventory management**
- ❖ **Content management:**
  - storing content and metadata of web sites
  - storing users' comments

see also <https://docs.mongodb.com/ecosystem/use-cases/>

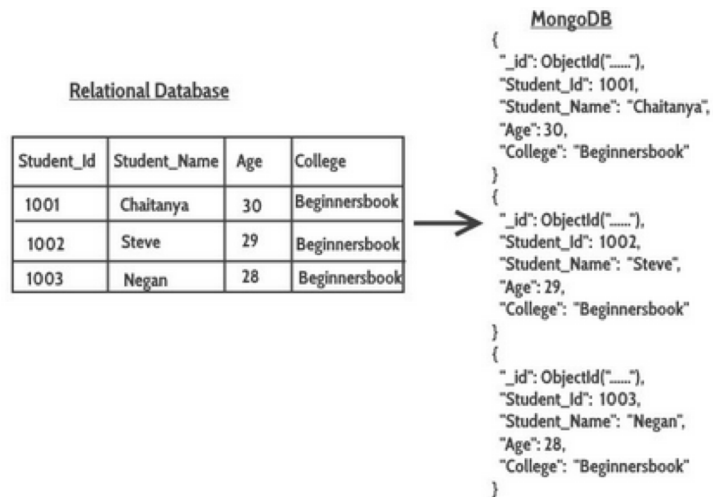
33

## JSON (JavaScript Object Notation),

- ❖ **Syntax (cf. [www.json.org](http://www.json.org)):**
  - JSON object = collection of name/value pairs in { }
  - the name is a string
  - the value can be
    - of a simple type: a number, string, true/false
    - an array: = sequence of values in [ ]
    - a nested JSON object: in { }
- ❖ **Remark:**
  - Two fields with the same name in principle allowed (in particular, also in BSON used by MongoDB)
  - most MongoDB interfaces disallow duplicate names because of the used representation (e.g., hash table)

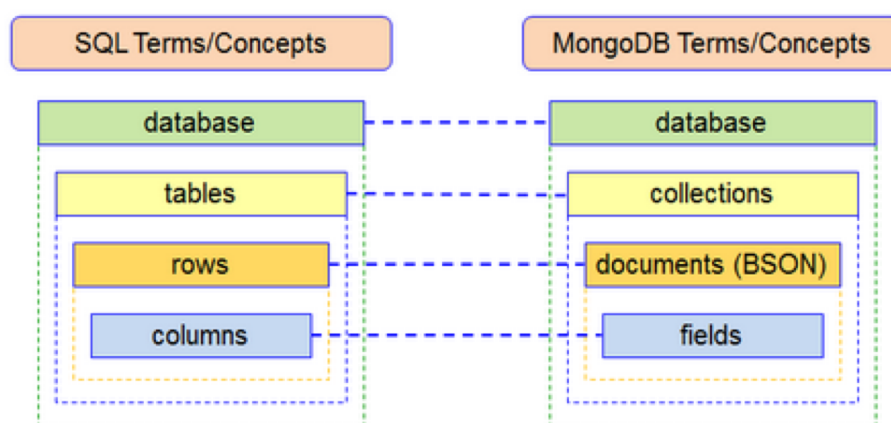
34

## Structure of the Database: SQL vs. MongoDB



35

## Structure of the Database: SQL vs. MongoDB (cont.)



36

## Data Modelling

### ❖ Basic principles:

- Keep intended data access in mind
- aim for atomic reads/writes if possible
- various ways to model 1:n and n:m relationships (choose the most suitable one for expected workload)
- denormalization is an option
- respect the max. document size of 16 MB
- use document validation (e.g.: warnings if a document gets an unexpected field: check if this is a bug or a feature of the schema design)
- decide which indexes shall be created

37

## Modelling 1:n Relationships

### ❖ Relational model:

- typically 2 tables (one for each of the entities)
- model 1:n relationship with a foreign key

### ❖ MongoDB: several possibilities:

- two collections with reference (analogous to a foreign key); reference can be `_id` or some key attribute: recommended for "one-to-millions"
- array of references in the "one-side" document, e.g.: array with `_ids` or order-numbers in customer document; recommended for "one-to-thousands"
- embed an array of documents in the other document, e.g.: array of order-documents in a customer document; for "one-to-few"

38

## Modelling 1:n Relationships (cont.)

### ❖ Combination of two referencing methods:

- useful if fast access from "both sides" is required

```

db.person.findOne()
{
  _id: ObjectID("AAF1"),
  name: "Kate Monster",
  tasks [
    ObjectID("ADF9"),
    ObjectID("AE02"),
    ObjectID("AE73")
    // etc
  ]
}

db.tasks.findOne()
{
  _id: ObjectID("ADF9"),
  description: "Write lesson plan",
  due_date: ISODate("2014-04-01"),
  owner: ObjectID("AAF1")
  // Reference to Person document
}

```

39

## Modelling n:m Relationships

### ❖ Relational model:

- typically three tables: auxiliary table for the relationship, one table for each entity,
- foreign keys from auxiliary table into each of the others

### ❖ MongoDB: several possibilities, e.g.:

- three collections with references as in relational case
- references as for 1:n relationship, but from both sides, e.g.: store references to courses in student document (i.e., array of course IDs) and references to students in course document (i.e., array of student IDs).

40

## Denormalization

### ❖ Idea:

- allow redundancy to avoid joins when querying
- disadvantage: updates are not atomic anymore
- makes most sense for high read-to-write ratio

### ❖ Example 1: denormalizing into "1-side"

```
db.products.findOne()
{
  name : 'left-handed smoke shifter',
  catalog_number: 1234,
  parts : [
    { id : ObjectID('AAAA'), name : '#4 grommet' },
    { id : ObjectID('F17C'), name : 'fan blade assembly' },
    { id : ObjectID('D2AA'), name : 'power switch' },
    // etc
  ]
}
```

41

## Denormalization (cont.)

### ❖ Example 2: denormalizing into "many-side"

- increases the cost of updates yet further
- high read-to-write ratio even more important

```
db.parts.findOne()
{
  _id : ObjectID('AAAA'),
  partno : '123-aff-456',
  name : '#4 grommet',
  product_name : 'left-handed smoke shifter',
  product_catalog_number: 1234,
  qty: 94,
  cost: 0.94,
  price: 3.99
}
```

42

## Document Validation

❖ MongoDB allows to specify rules for validating documents:

- defined for a collection
- in the form of a JSON Schema or a query
- allows specification of action: error/warning

```
db.createCollection( "contacts", {
  validator: { $jsonSchema: {
    bsonType: "object",
    required: [ "phone", "name" ],
    properties: {
      phone: {
        bsonType: "string",
        description: "string; required"
      },
      name: {
        bsonType: "string",
        description: "string; required"
      }
    }
  } },
  validationAction: "warn"
} )
```

43

## Replication

MongoDB uses **Master-Slave** distribution:

- MongoDB allows the definition of **replica sets**, i.e.: one **primary** (master) and several secondaries (slaves)
- Writes are only executed at the master and then propagated to the slaves.
- The secondary (slave) servers are in **hot standby**, i.e. they keep an up-to-date copy of the data and are ready to take over when the primary server goes down.
- When the secondary servers do not hear from the primary for more than 10 seconds (configurable), they hold an **election to vote for a new primary**.
- A replica set may also contain an **arbiter**, i.e.: a node that has a vote in an election but carries no data.

44

## Read Preferences

- ❖ MongoDB allows definition of Read preferences to define the read behaviour of a replica set:
  - **primary** resp. **secondary**: reads are *exclusively* executed on primary resp. secondary server;
  - **primaryPreferred** resp. **secondaryPreferred**: reads are executed on primary resp. secondary *if available*;
  - **nearest**: reads are executed on server which is nearest in terms of network latency.
  - For a transaction with reads & writes: read preference primary required
- ❖ **Example:** `db.getMongo().setReadPref('secondaryPreferred');`

45

## Write Concerns

- ❖ **Write concerns:**
  - A write operation is acknowledged in a replica set once the write is acknowledged by the primary server;
  - This behavior can be changed by write concerns on different levels (single operation, transaction, session)
  - Values, in particular "majority" or number of servers ("majority" =  $\lfloor N/2 \rfloor + 1$  with  $N = \text{\#servers in replica set}$ )
- ❖ **Example:** `db.books.insert(
 { name: "Mastering MongoDB 4.x", isbn: "1001" },
 { writeConcern: {w: 2, wtimeout: 5000} } )`  
 write is acknowledged when it is confirmed by 2 servers (i.e., primary plus one secondary). If the timeout of 5000 ms is exceeded, an error is returned.

46

## Read Concerns

### ❖ Read concerns:

- Can be defined for read operations (queries), e.g.: `find()`, `aggregate()`, `count()`, etc.;
- Values, in particular: "local": current value on the read target, "majority" (most recent value on a majority of the servers).
- default: "local".

❖ **Example:** `db.persons.find( { age: 22 } ).`  
`readConcern("majority").maxTimeMS(5000)`

read the most recent values from a majority of servers;  
 returns an error, if the timeout of 5000 ms is exceeded.

47

## Transactions

- ❖ **Multi-document transactions with ACID properties** have been introduced in version 4.0 (2018).
- ❖ **Atomicity of single document access** is guaranteed also without transactions (and recommended for performance).

### ❖ Example:

```
session = db.getMongo().startSession( );
session.startTransaction( { writeConcern: { w: "majority" } } );
db.persons.insertMany( [{ name: "alice", age: 22, location: "vienna"},
                        { name: "bob", age: 23, location: "graz"}]);
session.commitTransaction(); // respectively: session.abortTransaction();
session.endSession();
```

48



## Sharding

- [Sharding](#) is a method for distributing data across multiple machines. MongoDB uses sharding to support deployments with very large data sets and high throughput operations.
- Database systems with large data sets or high throughput applications can challenge the capacity of a single server. For example, high query rates can exhaust the CPU capacity of the server. Working set sizes larger than the system's RAM stress the I/O capacity of disk drives.

## Two methods for addressing system growth

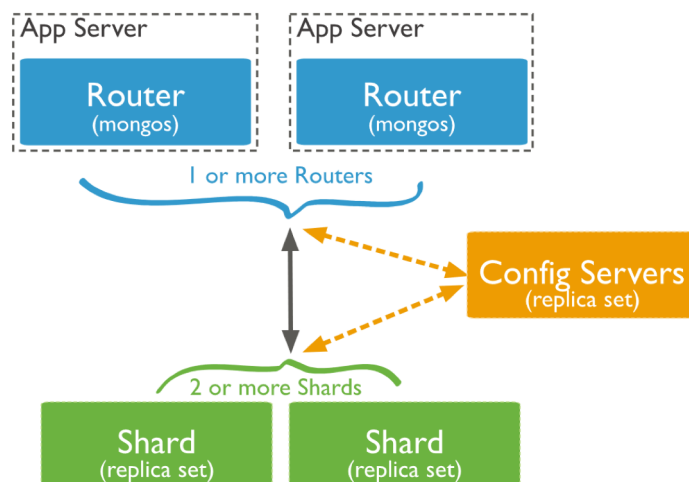
- **Vertical Scaling** involves increasing the capacity of a single server, such as using a more powerful CPU, adding more RAM, or increasing the amount of storage space.
- **Horizontal Scaling** involves dividing the system dataset and load over multiple servers, adding additional servers to increase capacity as required.

## Sharding

- ❖ MongoDB supports sharding on collection level (for horizontal scaling)
- ❖ A **sharded cluster** consists of the following elements:
  - Two or more shards, which can be replica sets
  - One or more query routers ("mongos")
  - A replica set of config servers (for metadata and configuration data of the entire cluster).
- ❖ Sharding is defined via a **shard key** (= field that exists in every document in the collection to be sharded)
- ❖ **Sharding Strategies:**
  - Hashed sharding: chunks defined by hash values
  - Ranged sharding: chunks as ranges of shard key values

51

## Components in a Sharded Cluster



52

## Advantages of Sharding

- **Reads / Writes**  
MongoDB **distributes the read and write workload** across the shards in the sharded cluster. For **queries that include the shard key** or the prefix of a compound shard key, mongos can **target the query at a specific shard** or set of shards.
- **Storage Capacity**  
Sharding **distributes data across the shards** in the cluster, allowing each shard to contain a subset of the total cluster data.
- **High Availability**  
A sharded cluster can **continue to perform partial read / write** operations even if one or more shards are unavailable. You can deploy config servers as **replica sets**

## Contents

- ❖ Motivation
- ❖ Four Categories of NoSQL Systems
- ❖ Further Aspects of Data Modelling
- ❖ Document Stores
- ❖ **MongoDB Primer**

## Basic MongoDB Commands (1)

| To do this                                    | Run this command    | Example            |
|---|---------------------|--------------------|
| Start server                                  | mongod              | mongod             |
| Start shell                                   | mongo               | mongo              |
| Show current database                         | db                  | db                 |
| Select or switch database <a href="#">[1]</a> | use <database name> | use mydb           |
| Execute a JavaScript file                     | load(<filename>)    | load (myscript.js) |
| Display help                                  | help                | help               |
| show all databases                            | show dbs            | show dbs           |
| Show all collections in current database      | show collections    | show collections   |

55

## Basic MongoDB Commands (2)

| To do this  | Run this command   | Example  |
|---|--|--|
| Insert new document in a collection <a href="#">[1]</a> | db.collection.insert(<document> )  | db.books.insert({"isbn": 9780, "title": "Alice in Wonderland", ...})                                       |
| Insert multiple documents into a collection             | db.collection.insertMany([<document1>, <document2>, ...])<br>-or-<br>db.collection.insert([<document1>, <document2>, ...]) | db.books.insertMany([{"isbn": 9780, "title": "Alice in Wonderland", ...}, {"isbn": 9781, "title": "..."}]) |
| Show all documents in the collection                    | db.collection.find()   | db.books.find()  |
| Filter documents by field value condition               | db.collection.find(<query>)  | db.books.find({"title": "Alice in Wonderland"})  |

56

## Basic MongoDB Commands (3)

| To do this   | Run this command  | Example  |
|--|---|--|
| Show only some fields of matching documents                    | <code>db.collection.find(&lt;query&gt;, &lt;projection&gt;)</code>    | <code>db.books.find({"title":"Alice in Wonderland"}, {title:true, _id:false})</code>           |
| Show first document that matches the query condition           | <code>db.collection.findOne(&lt;query&gt;, &lt;projection&gt;)</code> | <code>db.books.findOne({}, {_id:false})</code>   |
| Update specific fields of a single document                    | <code>db.collection.update(&lt;query&gt;, &lt;update&gt; )</code>     | <code>db.books.update({title : Alice in Wonderland"}, {\$set : {category : "Fiction"}})</code> |
| Remove certain fields of a single document the query condition | <code>db.collection.update(&lt;query&gt;, &lt;update&gt;)</code>      | <code>db.books.update({title : "Alice in Wonderland"}, {\$unset : {author:""}})</code>         |

57

## Basic MongoDB Commands (4)

| To do this  | Run this command  | Example   |
|---|---|---|
| Remove certain fields of all documents that match the query condition | <code>db.collection.update(&lt;query&gt;, &lt;update&gt;, {multi:true} )</code>                           | <code>db.books.update({title : "Alice in Wonderland"}, {\$unset : {author:""}}, {multi:true} )</code> |
| Delete all documents matching a query condition                       | <code>db.collection.remove(&lt;query&gt;)</code>  | <code>db.books.remove({"category" : "Fiction"})</code>  |
| Delete all documents in a collection                                  | <code>db.collection.remove({})</code>   | <code>db.books.remove({})</code>  |
| Create an index   | <code>db.collection.createIndex({indexField:type} )</code><br>Type 1 means ascending; -1 means descending | <code>db.books.createIndex({title:1} )</code>   |

58

## Basic MongoDB Commands (5)

| To do this  | Run this command  | Example   |
|---|---|---|
| Create an index on multiple fields (compound index) | <code>db.collection.createIndex({indexField1:type1, indexField2:type2, ...})</code> | <code>db.books.createIndex({title:1, author:-1})</code> |
| Show all indexes in a collection                    | <code>db.collection.getIndexes()</code>   | <code>db.books.getIndexes()</code>                      |
| Drop an index                                       | <code>db.collection.dropIndex({indexField:type} )</code>                            | <code>db.books.dropIndex({author:-1})</code>            |
| number of documents in the collection               | <code>cursor.count()</code>   | <code>db.books.find().count()</code>                    |
|   |   |   |
| Display formatted result                            | <code>cursor.pretty()</code>  | <code>db.books.find({}).pretty()</code>                 |

59

## Next Class

- Summarizing lecture content